

Bracha-Toueg Deadlock Detection Algorithm

- A process that suspects it is deadlocked, initiates a (Lai-Yang) snapshot to compute the wait-for graph.
- Each node u takes a local snapshot of:
 - requests it sent or received that weren't yet granted or dismissed;
 - and grant and dismiss messages in edges.
- Then it computes:
 - Out_u : the nodes it sent a request to (not granted)
 - In_u : the nodes it received a request from (not dismissed)
 - $Requests_u$: the number of grants u requires to become unblocked.

Bracha-Toueg Deadlock Detection Algorithm

Initially $notified_u = false$ and $free_u = false$ at all nodes u .

The **initiator** starts a deadlock detection run by executing *Notify*.

Notify_u: $notified_u \leftarrow true$
for all $w \in Out_u$ send NOTIFY to w
if $requests_u = 0$ then *Grant_u*
for all $w \in Out_u$ await DONE from w

Grant_u: $free_u \leftarrow true$
for all $w \in In_u$ send GRANT to w
for all $w \in In_u$ await ACK from w

While a node is awaiting DONE or ACK messages,
it can also process incoming NOTIFY and GRANT messages.

Bracha-Toueg Deadlock Detection Algorithm

Let u receive NOTIFY.

If $notified_u = false$, then u executes $Notify_u$.

u sends back DONE.

Let u receive GRANT.



If $requests_u > 0$, then $requests_u \leftarrow requests_u - 1$;



if $requests_u$ becomes 0, then u executes $Grant_u$.



u sends back ACK.

When the **initiator** has received DONE from all nodes in its *Out* set,
it checks the value of its *free* field.

If it is still *false*, the initiator concludes it is deadlocked.







Bracha-Toueg Deadlock Detection Algorithm

The Bracha-Toueg algorithm is **deadlock-free**:



The initiator eventually receives DONE's from all nodes in its *Out* set.

At that moment the Bracha-Toueg algorithm has terminated.

Two types of trees are constructed, similar to the echo algorithm:

1. NOTIFY/DONE's construct a tree T rooted in the initiator.
2. GRANT/ACK's construct disjoint trees T_v ,
rooted in a node v where from the start $requests_v = 0$.

The NOTIFY/DONE's only complete when all GRANT/ACK's have completed.



Time Complexity Analysis

- Message Complexity : 4 messages per node, <notify>, <grant>, <done>, <ack>.
- Time complexity: $2d$ hops, where d = diameter of the WFG.
 - One diameter to send <notify>/<grant> to the farthest node in the graph.
 - Another diameter to receive <done>/<ack> from the same.